



LANGUAGE MODELING FOR SPEECH RECOGNITION

HUNG-YI LEE 李宏毅

Why Language Model?

- Language model (LM): Estimated the probability of token sequence
 - Token sequence: $Y = y_1, y_2, \dots, y_n$
 - $P(y_1, y_2, \dots, y_n)$

HMM $Y^* = \arg \max_Y P(X|Y)P(Y)$


LM is usually helpful when your model outputs text

LAS $Y^* = \arg \max_Y \underline{P(Y|X)} \underline{P(Y)}$

Need paired data

Easy to collect

Why we need LM?

$$Y^* = \underset{Y}{\operatorname{arg\,max}} \underbrace{P(Y|X)} \underbrace{P(Y)}$$


Need paired data

Easy to collect

Words in Transcribed Audio

12,500 hours transcribed audio

= 12,500 x 60 x 130 ≈ 一億!


(哈利波特全套約 100 萬個詞)



Moschitta had been credited in The Guinness Book of World Records as the World's **Fastest Talker**

Source of video: <https://youtu.be/ExKCcndqK5c>

Why we need LM?

$$Y^* = \underset{Y}{arg \max} \underbrace{P(Y|X)} \underbrace{P(Y)}$$


Need paired data

Easy to collect

Words in Transcribed Audio

12,500 hours transcribed audio

= 12,500 x 60 x 130 ≈ 一億!

(哈利波特全套約 100 萬個詞)

BERT:

<https://youtu.be/UYPa347-DdE>

Just Words ...

BERT (一個巨大的 LM) 用了
30 億個以上的詞

N-gram

$P(\text{"wreck a nice beach"})$
 $= P(\text{wreck} | \text{START}) P(a | \text{wreck})$
 $P(\text{nice} | a) P(\text{beach} | \text{nice})$

- How to estimate $P(y_1, y_2, \dots, y_n)$
- Collect a large amount of text data as training data
 - However, the token sequence y_1, y_2, \dots, y_n may not appear in the training data
- *N-gram language model*: $P(y_1, y_2, \dots, y_n) = P(y_1 | \text{BOS}) P(y_2 | y_1) \dots P(y_n | y_{n-1})$
 - E.g. Estimate $P(\text{beach} | \text{nice})$ from training data

← 2-gram

$$P(\text{beach} | \text{nice}) = \frac{C(\text{nice beach})}{C(\text{nice})}$$

Count of "nice beach" ←

Count of "nice" ←

- It is easy to generalize to 3-gram, 4-gram

Challenge of N-gram

- The estimated probability is not accurate.
 - Especially when we consider n-gram with large n
 - Because of data sparsity (many n-grams never appear in training data)

Training Data:



The dog ran
The cat jumped

$$P(\text{jumped} \mid \text{the, dog}) = \cancel{0} \quad 0.0001$$

$$P(\text{ran} \mid \text{the, cat}) = \cancel{0} \quad 0.0001$$

Give some small probability

This is called **language model smoothing**.

Continuous LM

- Recommendation system

				
A	5	5		1
B	5	5?	1	
C	5	5		2
D	1		4	4
E		1	5	4

Matrix Factorization

Ref: https://youtu.be/iwh5o_M4BNU?t=4673

Borrowing the idea from recommendation system

Continuous LM

Vocabulary

		dog h^1	cat h^2	child
{	ran v^1	2 n_{11}	3		1
	jumped v^2	0	2 n_{22}		1
	cried v^3	0	0		3
	laughed v^4	0	0		3
				

history

Not observed

Count of "cat jumped"

v^i, h^j are vectors to be learned

$$n_{12} = v^1 \cdot h^2$$
$$n_{21} = v^2 \cdot h^1 \dots$$

Minimizing

$$L = \sum_{(i,j)} (v^i \cdot h^j - n_{ij})^2$$

v^i, h^j found by gradient descent

Borrowing the idea from recommendation system

Continuous LM

Vocabulary

		dog h^1	cat h^2	child
{	ran v^1	2 n_{11}	3		1
	jumped v^2	0	2 n_{22}		1
	cried v^3	0	0		3
	laughed v^4	0	0		3
				

Not observed

Count of “cat jumped”

history

History “dog” and “cat” can have similar vector h^{dog} and h^{cat}

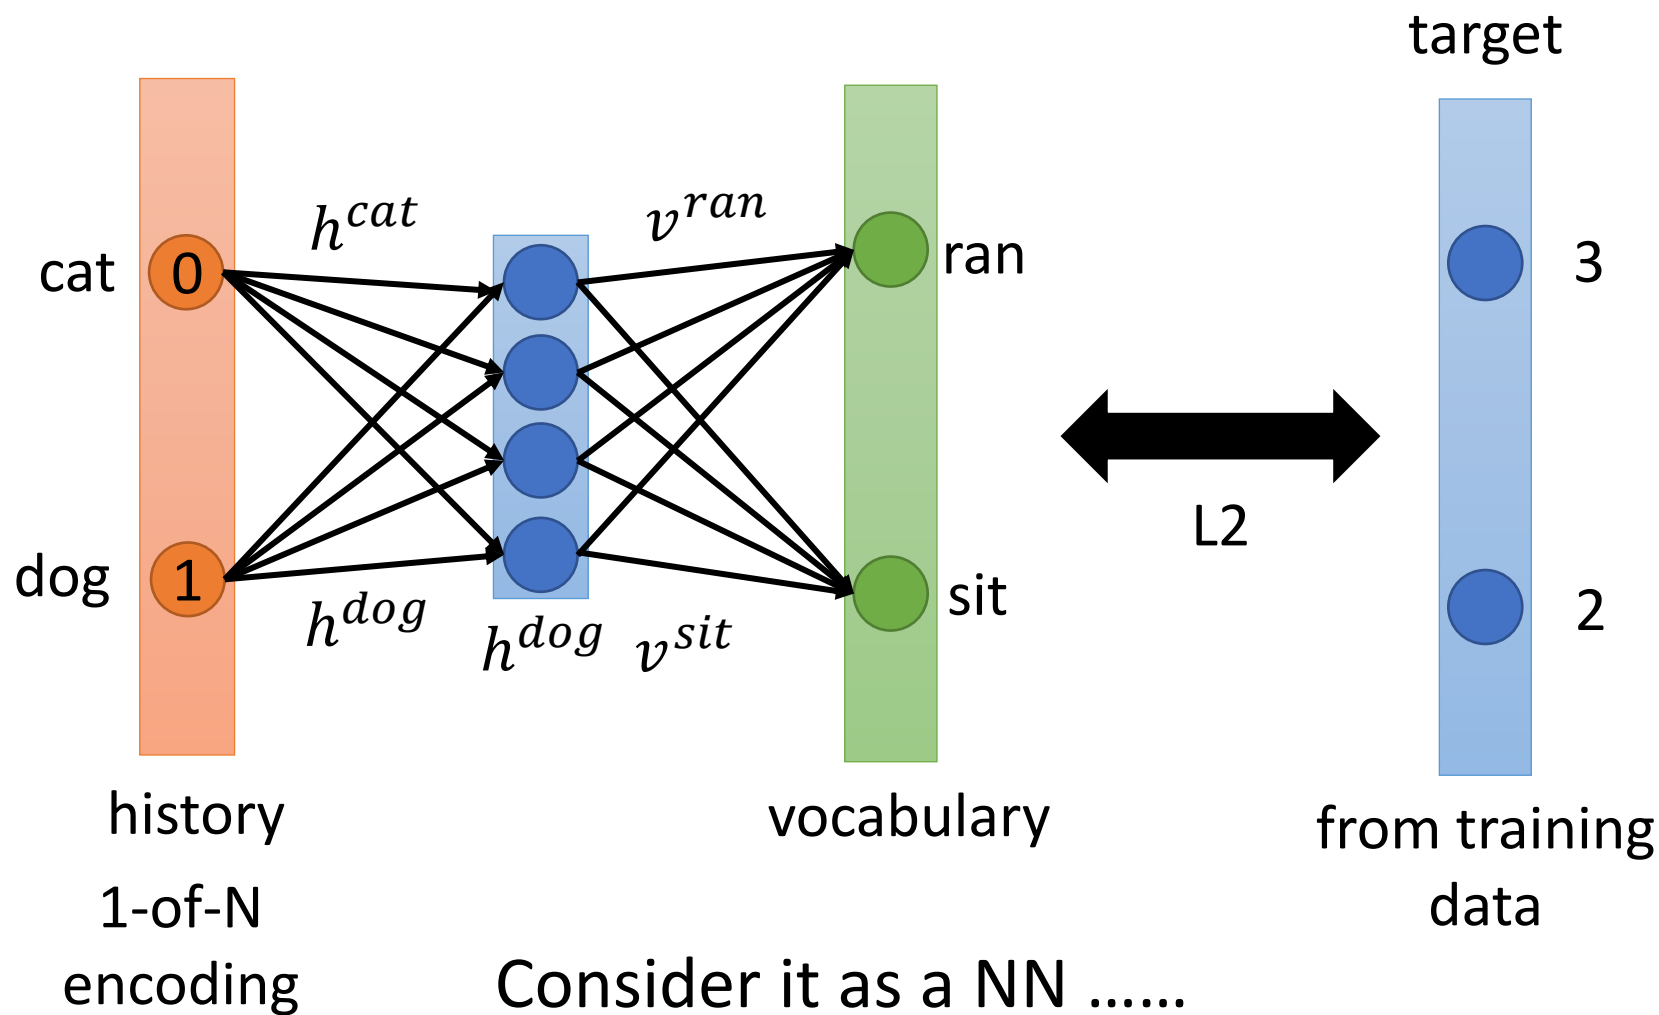
If $v^{\text{jumped}} \cdot h^{\text{cat}}$ is large, $v^{\text{jumped}} \cdot h^{\text{dog}}$ would be large accordingly.

Even if we have never seen “dog jumped ...”

Smoothing is automatically done.

Continuous LM

$$L = \sum_{(i,j)} (v^i \cdot h^j - n_{ij})^2$$




NN-based LM

- Training:

Collect data:

潮水 退了 就 知道 誰 ...
不爽 不要 買 ...
公道價 八萬 一 ...
.....

**Learn to predict
the next word**

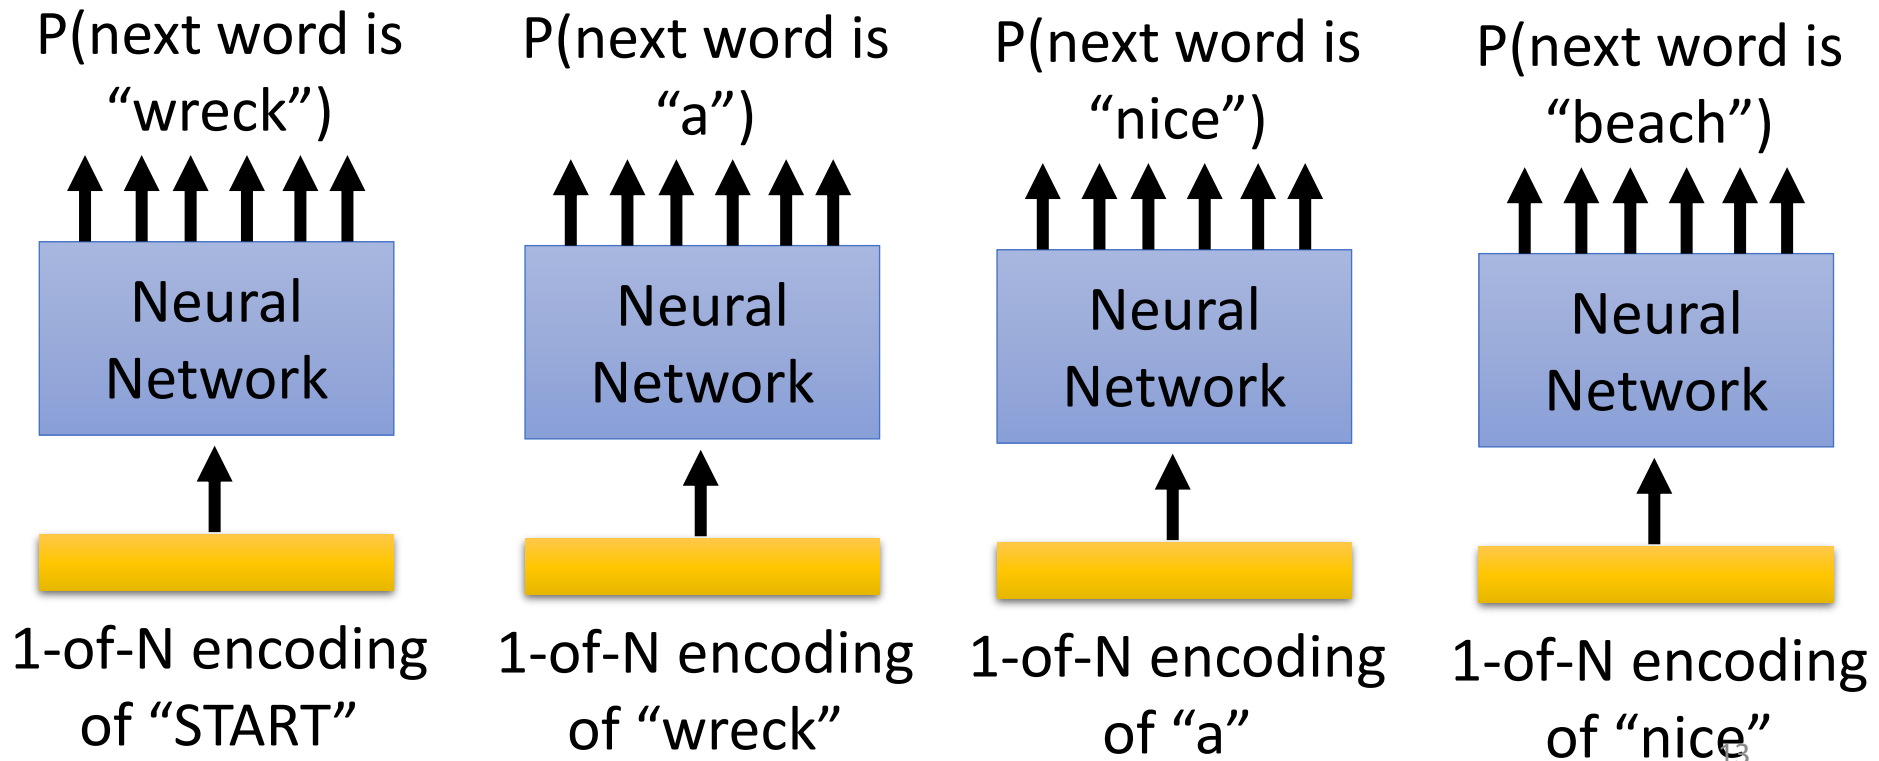



NN-based LM

$P(\text{"wreck a nice beach"})$

$=P(\text{wreck} | \text{START})P(a | \text{wreck})P(\text{nice} | a)P(\text{beach} | \text{nice})$

$P(b | a)$: the probability of NN predicting the next word.

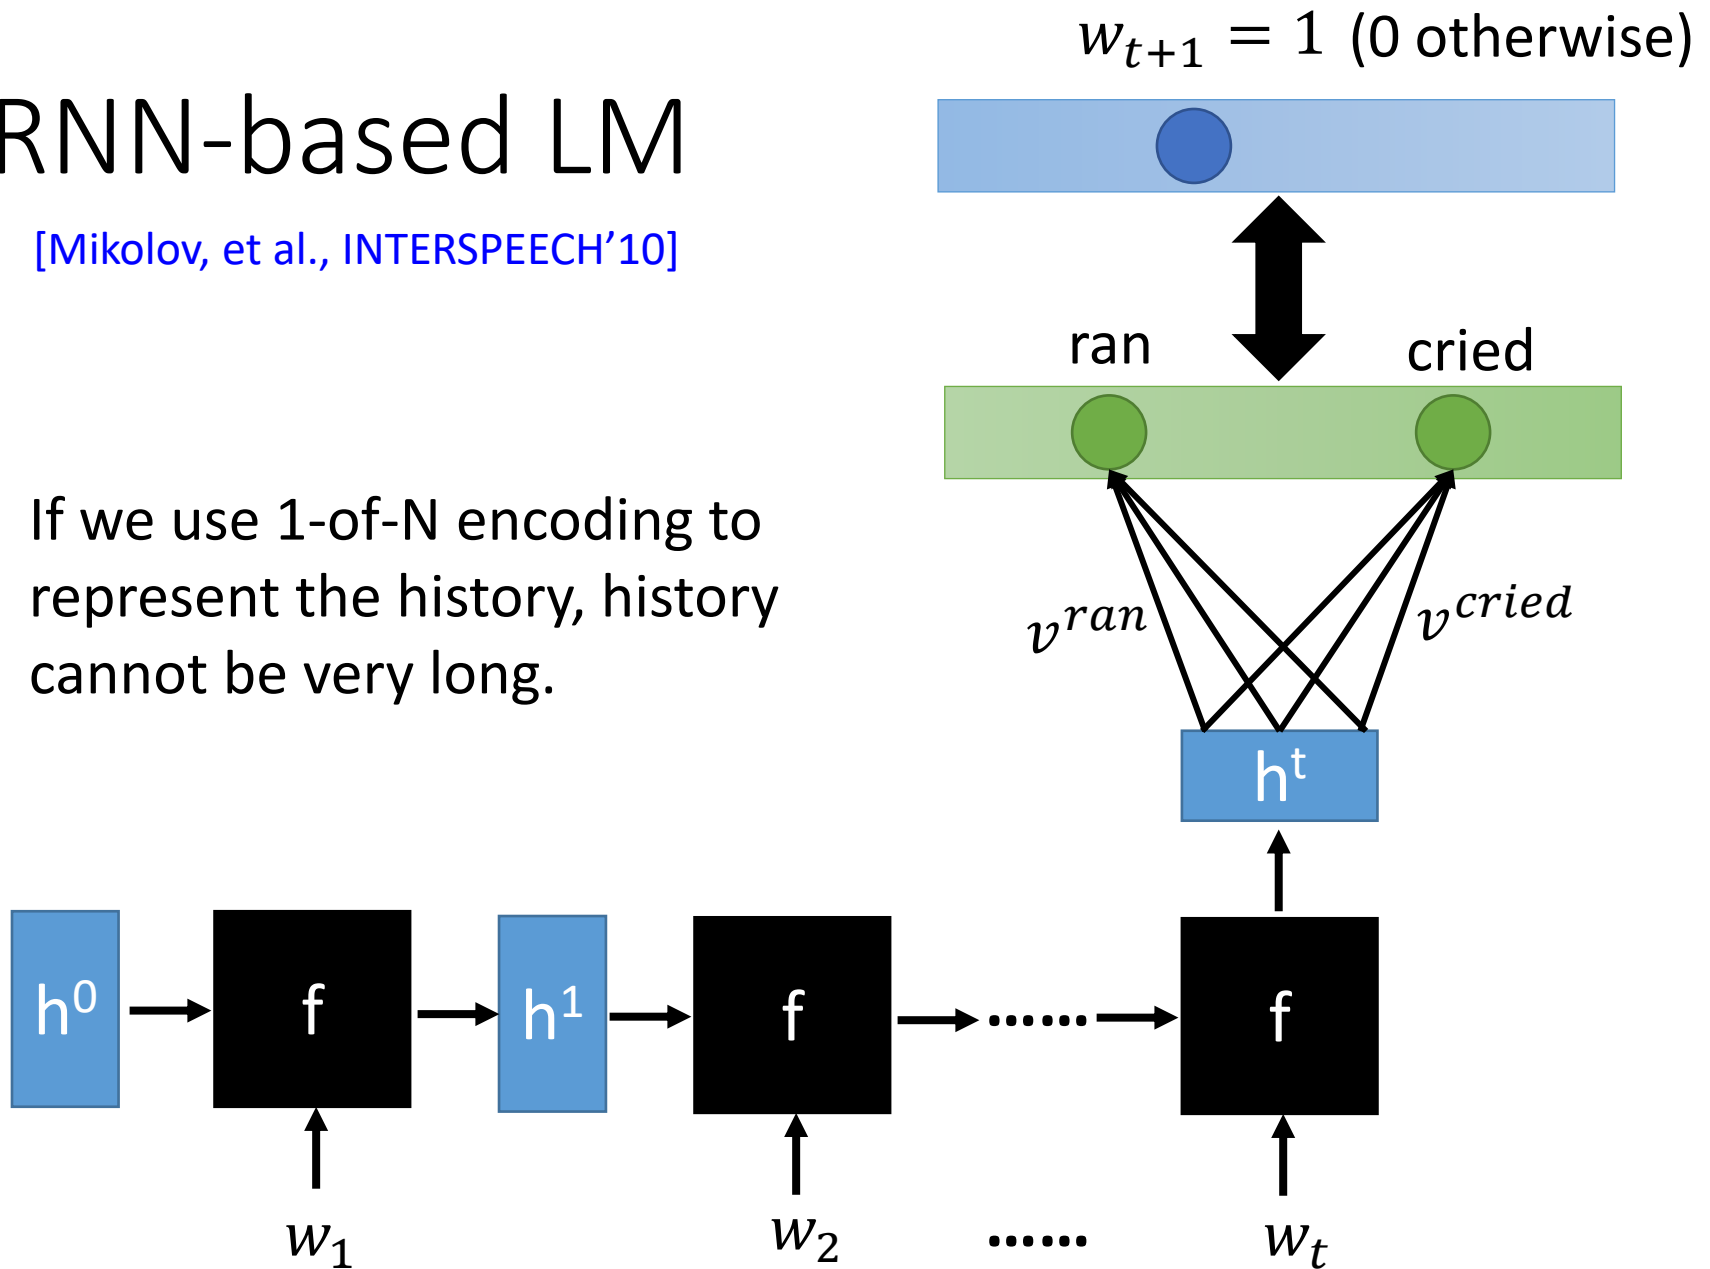





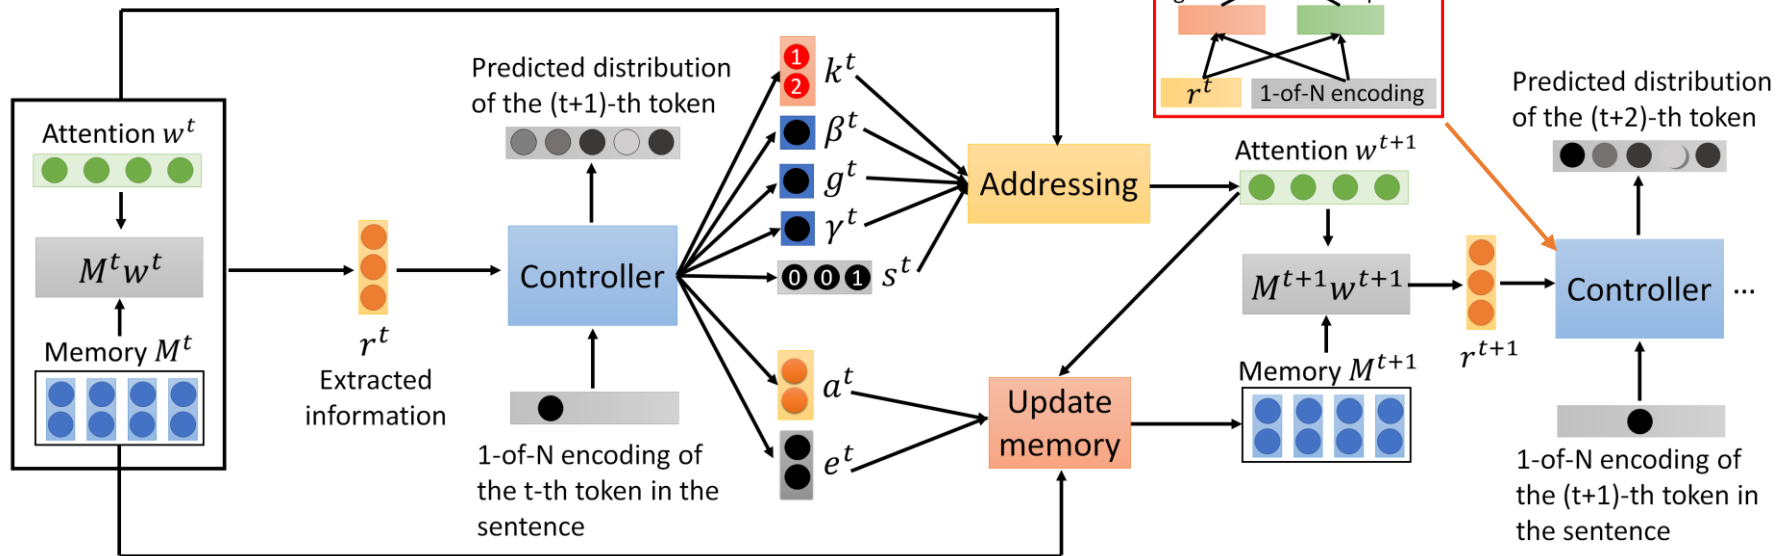
RNN-based LM

[Mikolov, et al., INTERSPEECH'10]

If we use 1-of-N encoding to represent the history, history cannot be very long.



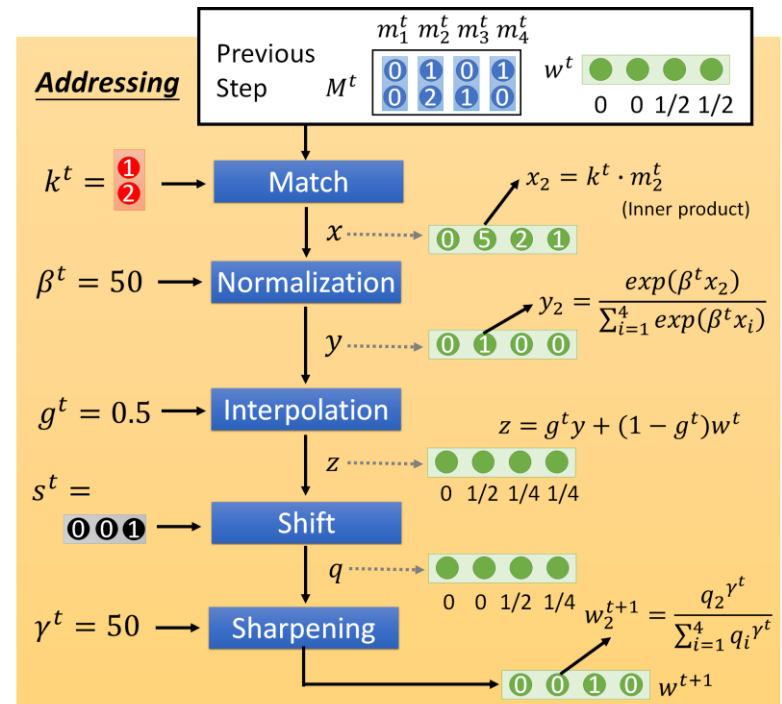
Can be very complex



[Ko, et al., ICASSP'17]

LSTM with proper optimization and regularization can be good.

[Merity, et al., ICLR'18]



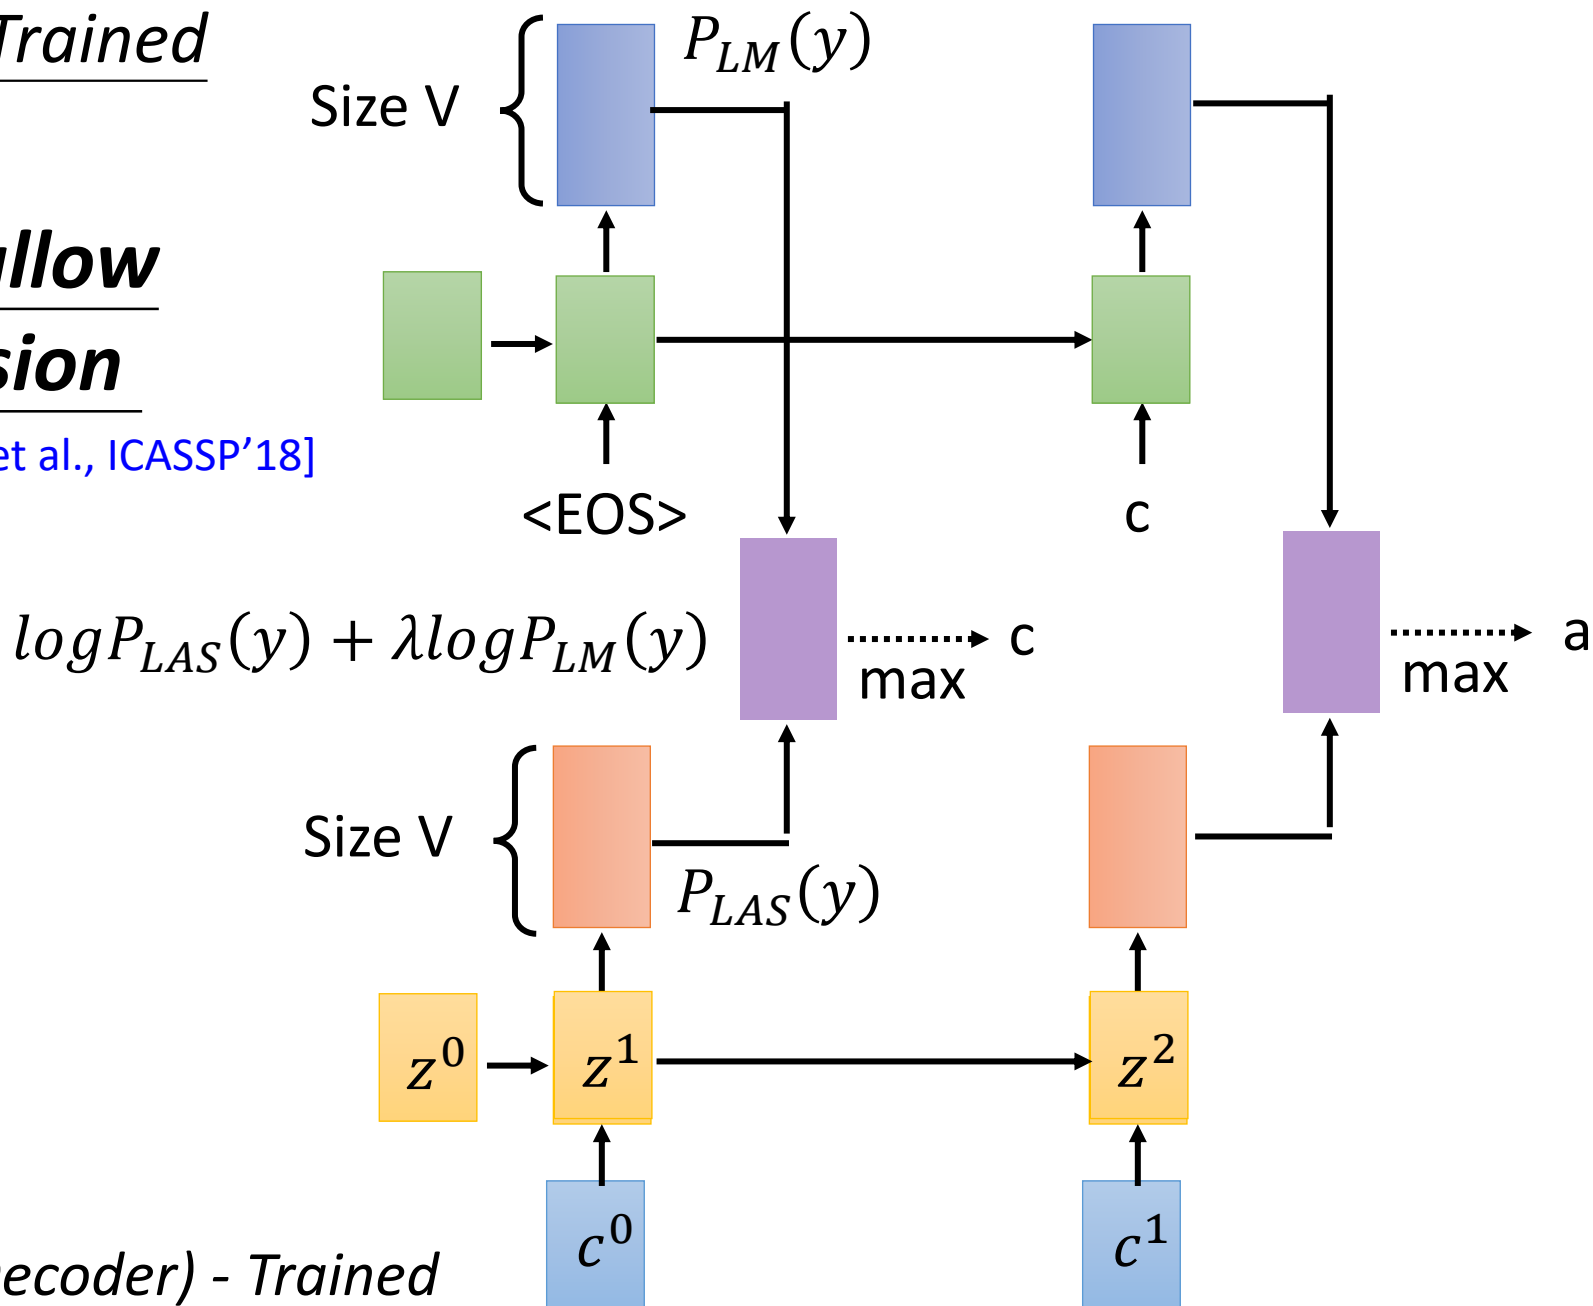
How to use LM to improve LAS?

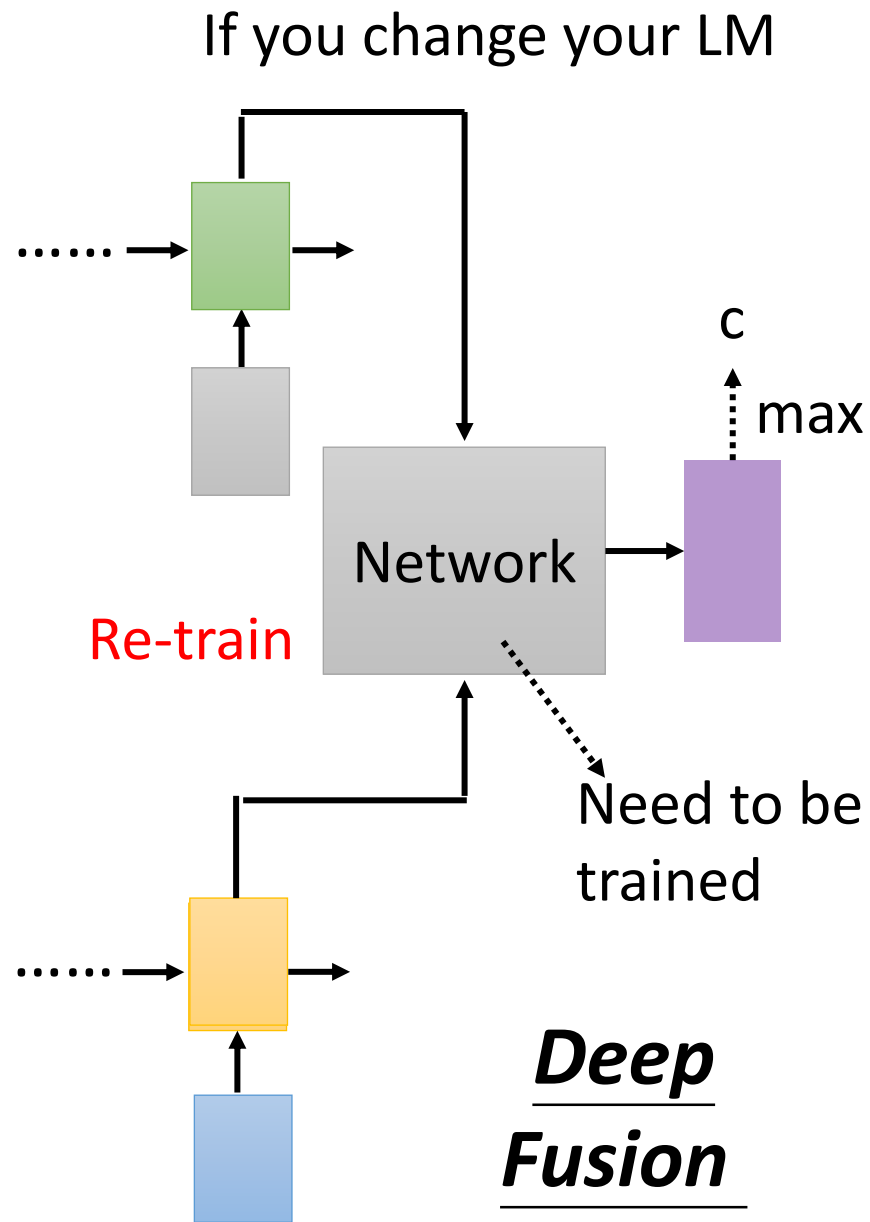

		how to integrate	
when to integrate		Output	Hidden
	After Training	Shallow Fusion	Deep Fusion
	Before Training		Cold Fusion

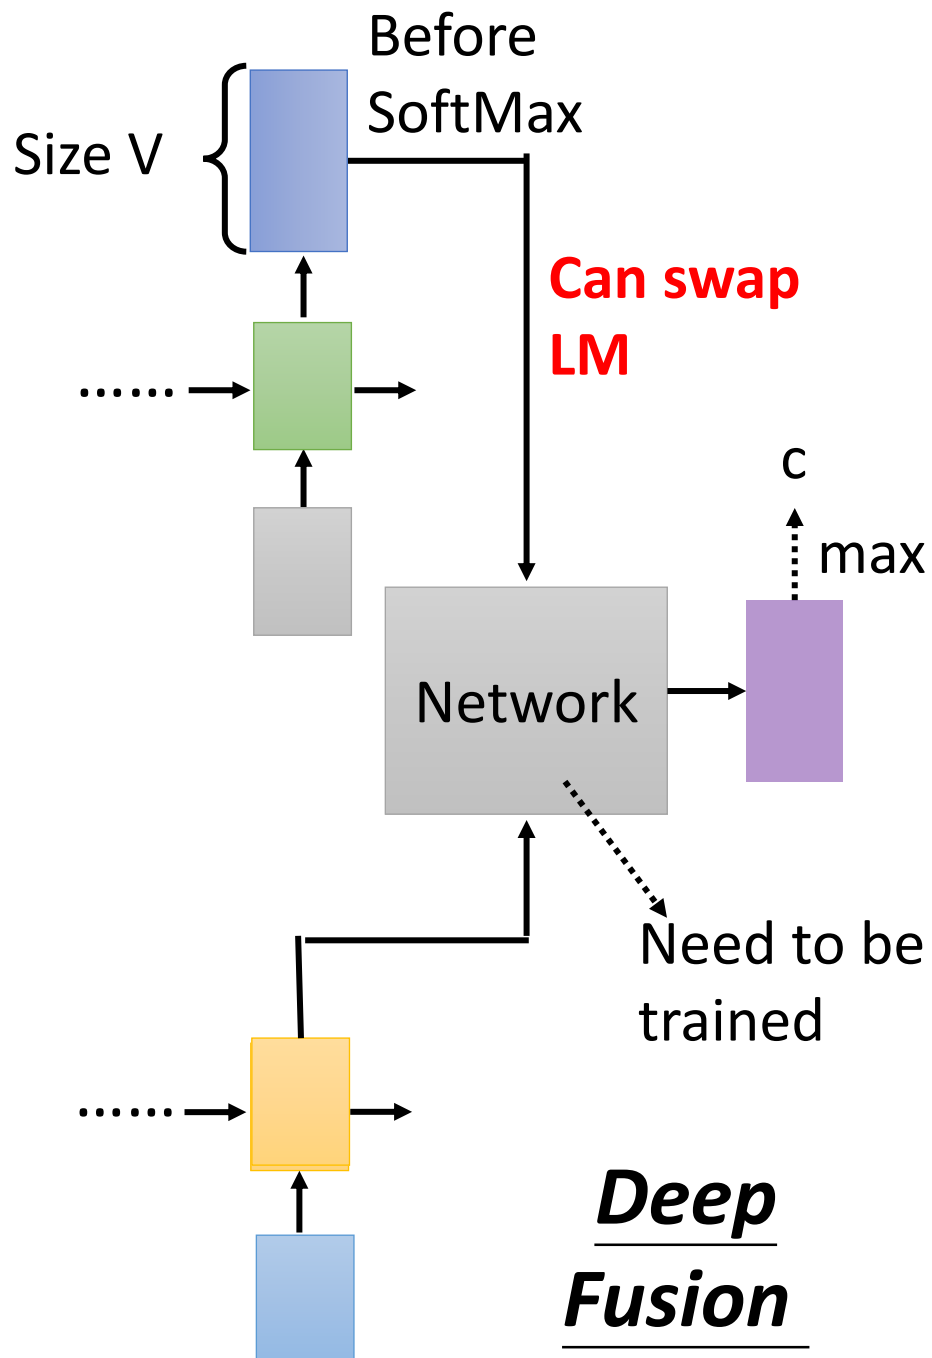

LM - Trained

**Shallow
Fusion**

[Kannan, et al., ICASSP'18]








Cold Fusion

[Sriram, et al., INTERSPEECH'18]

LM is already
trained

- LAS converges faster during training
- LAS has to be trained again if you have a new LM.

LAS is trained
from scratch



Concluding Remarks

		how to integrate	
when to integrate		Output	Hidden
	After Training	Shallow Fusion	Deep Fusion
	Before Training		Cold Fusion

Reference

- [Bengio, et al., JMLR'03] Yoshua Bengio, Réjean Ducharme, Pascal Vincent, Christian Janvin, A neural probabilistic language model, The Journal of Machine Learning Research, March 2003
- [Mikolov, et al., INTERSPEECH'10] Tomáš Mikolov, Martin Karafiát, Lukáš Burget, Jan Černocký, Sanjeev Khudanpur, Recurrent Neural Network Based Language Model INTERSPEECH, 2010
- [Ko, et al., ICASSP'17] Wei-Jen Ko, Bo-Hsiang Tseng, Hung-yi Lee, “Recurrent Neural Network based Language Modeling with Controllable External Memory”, ICASSP, 2017
- [Merity, et al., ICLR'18] Stephen Merity, Nitish Shirish Keskar, Richard Socher, Regularizing and optimizing LSTM language models, ICLR, 2018

Reference

- [Gulcehre, et al., arXiv'15] Caglar Gulcehre, Orhan Firat, Kelvin Xu, Kyunghyun Cho, Loic Barrault, Huei-Chi Lin, Fethi Bougares, Holger Schwenk, Yoshua Bengio, On Using Monolingual Corpora in Neural Machine Translation, arXiv, 2015
- [Sriram, et al., INTERSPEECH'18] Anuroop Sriram, Heewoo Jun, Sanjeev Satheesh, Adam Coates, Cold Fusion: Training Seq2Seq Models Together with Language Models, INTERSPEECH, 2018
- [Kannan, et al., ICASSP'18] Anjuli Kannan, Yonghui Wu, Patrick Nguyen, Tara N. Sainath, Zhifeng Chen, Rohit Prabhavalkar, An analysis of incorporating an external language model into a sequence-to-sequence model, ICASSP, 2018